

House Price Prediction Mini Project

Aim

To predict house prices based on various features using a linear regression model.

Algorithm

Linear Regression

Description of Workflow

1. Data Collection

Load the dataset containing features such as area, number of bedrooms, and other relevant attributes.

2. Data Preprocessing

Clean the data by handling missing values and encoding categorical variables if necessary.

Split the dataset into training and testing sets.

3. Model Training

Implement the linear regression algorithm using the training data.

Use gradient descent to optimize the model parameters (weights and bias).

4. Model Evaluation

Evaluate the model's performance using metrics such as Mean Squared Error (MSE) or R-squared on the testing set.

5. Prediction

Use the trained model to make predictions on new data.

Source Code (with Simulated Data and Fixes)

```
import numpy as np
import matplotlib.pyplot as plt

# Simulate a dataset
def load_data():
    np.random.seed(0) # For reproducibility
    # Simulating 100 data points
    x_train = np.random.rand(100, 1) * 1000 # Feature:
    Area in square feet
    y_train = 150 + 0.3 * x_train + (np.random.randn(100,
    1) * 50) # Price with some noise
    return x_train, y_train

# Compute cost function
def compute_cost(x, y, w, b):
    m = x.shape[0]
    total_cost = np.sum((x.dot(w) + b - y) ** 2) / (2 * m)
    return total_cost

# Compute gradient
def compute_gradient(x, y, w, b):
    m = x.shape[0]
    dj_dw = (1/m) * x.T.dot(x.dot(w) + b - y) # Transpose
    for correct shape
    dj_db = (1/m) * np.sum(x.dot(w) + b - y)
    return dj_dw, dj_db

# Gradient descent
def gradient_descent(x, y, w_in, b_in, alpha, num_iters):
    w = w_in
    b = b_in
    for i in range(num_iters):
        dj_dw, dj_db = compute_gradient(x, y, w, b)
        w -= alpha * dj_dw
        b -= alpha * dj_db
    return w, b

# Main execution
x_train, y_train = load_data()
initial_w = np.zeros((x_train.shape[1],1)) # Column vector
```


```
initial_b = 0
alpha = 0.01
iterations = 1500

w, b = gradient_descent(x_train, y_train, initial_w,
                        initial_b, alpha, iterations)
print("w, b found by gradient descent:", w.flatten(), b)

# Optional: Plotting the results
plt.scatter(x_train, y_train, color='blue', label='Data
points')
plt.plot(x_train, x_train.dot(w) + b, color='red',
label='Regression line')
plt.xlabel('Area (sq ft)')
plt.ylabel('Price')
plt.title('House Price Prediction')
plt.legend()
plt.show()
```

Output for the Code

w, b found by gradient descent: [1.16636235] -3.63029143940436



This output indicates the slope and intercept of the linear regression model, which can then be used to make predictions on house prices based on the input features.